



DyHuCoG: A Dynamic Hypergraph Cooperative Game for Preference-aware Recommendation

Mouad Louhichi^{1*}Redwane Nesmaoui¹Mohamed Lazaar¹¹National Higher School of Computer Science and Systems Analysis, Mohammed V University in Rabat, Morocco* Corresponding author's Email: mouad_louhichi@um5.ac.ma

Abstract: Recommender systems must balance accuracy, diversity, and robustness to enhance the user experience and engagement. We present DyHuCoG (Dynamic Hypergraph Cooperative Game), a cooperative-game-enhanced hypergraph recommender that models user–item–context interactions as dynamic coalitions. A preference-aware Shapley estimator (Monte Carlo) quantifies the marginal utility of each interaction, which is then injected as dynamic hyperedge weights into a lightweight hypergraph neural network with an interaction-level-attention gate. This design enables the model to adaptively weight user, item, and context signals, capturing nuanced preferences that conventional collaborative filtering and graph-based models often ignore. Evaluated on the MovieLens-1M and Amazon-Book datasets, DyHuCoG consistently outperforms strong baselines, including matrix factorization (MF), neural collaborative filtering (NCF), Light Graph Convolutional Network (LightGCN), Dual Contrastive Learning for Recommendation (RecDCL), Hypergraph Contrastive Collaborative Filtering (HCCF) and Hypergraph Projection Enhanced Collaborative Filtering (HPCF). Relative to HPCF (state-of-the-art baseline), DyHuCoG achieves improvements in Normalized Discounted Cumulative Gain at cutoff 20 (NDCG@20) and recall at cutoff 20 (Recall@20) by +9.77%/+12.58% (MovieLens-1M) and +13.33%/+16.16% (Amazon-Book) and further increases coverage and diversity. Although Shapley estimation introduces additional computation, end-to-end training remains practical at approximately 1.78× the duration of HPCF. These results demonstrate that integrating preference-aware cooperative game attributions directly into hypergraph message passing yields a lightweight, training-stable module that jointly improves ranking quality, coverage, and diversity for large-scale recommendation tasks.

Keywords: Recommender systems, Hypergraph neural networks, Cooperative games, Shapley value, Coverage, Diversity.

1. Introduction

Digital platforms increasingly rely on recommender systems to guide users toward relevant content, products and services. While accuracy has long been the primary objective of recommendation research, there is growing recognition that diversity is equally critical for enhancing user satisfaction, engagement, and long-term platform health[1]. Recommender systems that focus solely on accuracy often produce highly similar or redundant suggestions, leading to user fatigue, reduced engagement, and the so-called "filter bubble" effect where users are repeatedly exposed to a narrow subset of items[2].

Most collaborative filtering (CF) and graph-based recommenders infer preferences primarily from past user–item interactions[3]. Although effective, such models often behave like black boxes and treat edges homogeneously, allowing head-item signals to dominate propagation while down-weighting weak but informative evidence. As preferences and contexts shift, static weighting further amplifies popularity bias and exposes brittleness to sparsity, resulting in narrow recommendations[4].

We introduce DyHuCoG (Dynamic Hypergraph Cooperative Game) a hypergraph neural recommender that embeds preference-aware Shapley values into message-passing. For each user–item–context interaction, DyHuCoG estimates a contribution score, its Shapley value, via efficient

Monte Carlo sampling[5]. These scores are injected as dynamic hyperedge weights so that the information flow reflects the extent to which each interaction improves the underlying ranking objective. An interaction-level attention gate complements this mechanism by amplifying high-contribution signals and dampening noisy signals. Training uses a composite objective that combines a pairwise ranking loss with an explicit diversity regularizer to encourage broader exposure without sacrificing accuracy[6]. Beyond performance, the Shapley-based weights provide an interpretable lens on why items are recommended, helping to diagnose popularity effects and contextual sensitivities[7].

We evaluated DyHuCoG on two public benchmarks with contrasting characteristics: MovieLens-1M (denser, explicit feedback) and Amazon-Book (sparser, long-tail implicit feedback). Comparisons against representative baselines, Matrix Factorization (MF), Neural Collaborative Filtering (NCF), Light Graph Convolutional Network (LightGCN), Dual Contrastive Learning for Recommendation (RecDCL)[8], Hypergraph Contrastive Collaborative Filtering (HCCF), and Hypergraph Projection Enhanced Collaborative Filtering (HPCF)[9], use standard top-K metrics alongside catalog coverage and intra-list diversity. While Shapley estimation adds computation, sampled attribution, and a lightweight propagation backbone, it keeps training practical and amenable to large catalogs.

The remainder of this paper is organized as follows: Section 2 reviews related work on recommender systems and cooperative game theory. Section 3 presents the DyHuCoG framework and its integration with graph neural networks. Section 4 describes the datasets, metrics, and experiments, followed by the results section. Section 5 concludes the paper and suggests directions for future research.

2. Literature review

Recommender systems are fundamental to modern digital platforms; however, they face persistent challenges in balancing accuracy, diversity, and adaptability, especially as user preferences and item catalogs become increasingly dynamic. Early methods, such as matrix factorization (MF), provided foundational accuracy but struggled with data sparsity, limiting their effectiveness in real-world, large-scale environments[10]. In our experiments on MovieLens-1M and Amazon-Book, MF consistently underperformed on key ranking metrics, such as NDCG@20 and Recall@20, particularly in the sparser Amazon-Book dataset, highlighting its

limitations in handling sparse data and new users or items.

Neural collaborative filtering (NCF) has advanced the field by leveraging deep learning to capture complex nonlinear user-item relationships; however, this increased representational power often comes at the expense of interpretability and adaptability[11]. NCF achieves moderate improvements in NDCG@20 and Recall@20 over MF, yet remains a "black box" and is still hampered by the lack of diversity in recommendations. This is evident in our results, where the NCF gains are incremental and its performance on diversity metrics remains unsatisfactory.

The introduction of graph-based recommender systems, particularly LightGCN and HCCF, has marked a significant leap in this field. LightGCN, with its streamlined architecture, has become a state-of-the-art baseline owing to its efficiency and strong empirical performance, achieving the highest NDCG@20 and Recall@20 among traditional baselines in both datasets[12]. However, its reliance on historical user-item interactions reinforces popularity bias, which hinders the diversity of recommendations[13]. Our data show that, while LightGCN excels in ranking metrics, its coverage and diversity scores are limited. The HCCF extends this paradigm by modeling higher-order connectivity through neural message passing, capturing more intricate relationships, but at the cost of increased computational complexity and a heightened risk of overfitting. While both models excel in ranking accuracy, they primarily optimize metrics such as NDCG@20 and Recall@20 and do not explicitly address the intertwined challenges of diversity.

Recent hypergraph recommenders leverage higher-order relationships among users, items, and contexts through hyperedge-based propagation, which better captures multiway co-occurrence under sparsity. In parallel, contrastive graph collaborative filtering maximizes agreement across perturbed graph views (e.g., node/edge dropout and feature masking) and has become a strong state-of-the-art in sparse-feedback regimes. To reflect this progress and ensure a fair comparison, we included HCCF, a recent hypergraph CF baseline, two recent 2024–2025 hypergraph/contrastive methods, RecDCL and HPCF, in both the review and experimental comparisons[14].

In pursuit of greater transparency and adaptability, recent research has turned to game-theoretic frameworks, with particular attention to the Shapley value for principled contribution attribution[15]. Recent studies have demonstrated that integrating adaptive Shapley values into collaborative filtering

models can yield notable improvements in diversity and precision by quantifying the marginal impact of each user-item interaction[16]. These advances have enhanced the interpretability and fairness of recommendations; however, the computational demands of exact Shapley value calculations remain a barrier, and most implementations are static, limiting their ability to dynamically influence diversity in real-time. Moreover, these methods are rarely evaluated using standard ranking metrics, such as NDCG and Recall making direct comparisons with state-of-the-art graph-based models difficult.

Context-aware recommender systems have sought to overcome these limitations by combining collaborative, content-based, and knowledge-driven approaches, often incorporating deep neural networks and real-time context signals[17]. While such models can alleviate issues when rich metadata are available, they frequently introduce significant computational overhead and lack principled mechanisms for attributing the influence of context on recommendations[18]. Similarly, dedicated solutions for diversity, such as ontology-based algorithms [19], meta-learning, and hybrid diffusion processes[20], have shown promise but often face trade-offs with accuracy, depend heavily on extensive metadata, and struggle to scale efficiently[21]. Notably, most graph-based models do not natively support these enhancements and typically require auxiliary modules or loss terms to address the diversity challenges[22]. In our experiments, models that attempted to increase diversity often did so at the expense of NDCG@20 and Recall@20, underscoring the difficulty in simultaneously achieving all objectives.

Despite these advances, a unified framework that dynamically models user-item-context interactions, leverages principled contribution attribution, and inherently optimizes accuracy and diversity remains elusive[23]. The DyHuCoG framework was designed to bridge this critical gap by integrating dynamic cooperative game theory with lightweight hypergraph neural networks, offering a scalable, interpretable, and robust solution for next-generation recommendation systems. As demonstrated in our results, DyHuCoG achieves substantial improvements in NDCG@20 and Recall@20 over all baselines while also delivering higher coverage and diversity, confirming its effectiveness in addressing the core limitations of existing methods.

Taken together, prior methods improve top-K accuracy but rarely couple representation learning with contribution-aware propagation, which can counteract popularity dominance and sparsity. Our approach addresses this by integrating preference-

aware Shapley attributions into hypergraph message passing to dynamically weight interactions while learning with a diversity-aware objective. To align with current practice, we report NDCG and Recall@K, together with coverage and intra-list diversity; where appropriate, we apply paired significance tests and report runtime.

3. Methodology

3.1 Problem formulation and notation

We address the top-N recommendation problem in a user-item interaction setting that is optionally enriched with contextual information. Let $\mathcal{U} = \{u_1, \dots, u_{|\mathcal{U}|}\}$ denote the set of users, $\mathcal{I} = \{i_1, \dots, i_{|\mathcal{I}|}\}$ the set of items, and $\mathcal{C}$ the set of possible contexts (e.g., time, location, or session). The observed interactions are represented as a bipartite graph $\mathcal{G} = (\mathcal{U} \cup \mathcal{I}, \mathcal{E})$, where $(u, i) \in \mathcal{E}$ if user u has interacted with item i . Each interaction may be associated with a context vector $c_{u,i} \in \mathcal{C}$.

The objective is to learn a scoring function $f(u, i, c_{u,i})$ that ranks items for each user, maximizing both accuracy (as measured by NDCG@20, Recall@20, and diversity).

3.2 DyHuCoG framework overview

The DyHuCoG framework combines cooperative game modeling, Shapley value attribution, and hypergraph learning to generate adaptive recommendations. As shown in Fig. 1, the process begins with the input layer, where users, items, and context (e.g., genre) are encoded as nodes. These interactions form the basis for coalition scoring in the cooperative game module, which evaluates the utility in terms of accuracy, diversity, and contextual alignment. Next, the Shapley value network estimates the contribution of each entity using preference-aware Shapley values. These values serve as dynamic weights that guide the flow of information in the graph. The hypergraph Graph Neural Network (GNN) then propagates signals across higher-order connections, updating the node embeddings while respecting the Shapley-derived importance. To refine the predictions, an attention mechanism highlights the most salient user-item-context relationships. Finally, the output layer produces weighted recommendation scores. The training objective jointly optimizes ranking accuracy, diversity, context-awareness, and regularization. This modular design allows DyHuCoG to adapt to evolving preferences while ensuring the fair attribution of interactions.




Figure. 1 Workflow of the proposed DyHuCoG framework

3.3 Dynamic cooperative game modeling

We instantiate a dynamic cooperative game (N, v) whose players are the users, items, and contexts. The coalition value quantifies the quality of the recommendation configuration produced using only the entities in $S \subseteq N$. Formally, we measure the multi-objective utility as follows.

$$v(S) = \alpha \cdot \text{NDCG@20}(S) + \beta \cdot \text{Diversity}(S) + \gamma \cdot \text{ContextScore}(S) \quad (1)$$

The weights (α, β, γ) govern the trade-off among these objectives and are tuned for validation splits. To bias coalitions toward user-preference consistency, we add a preference-weighted variant that increases value when item choices agree with historical or content-based affinities:

$$v_{\text{pref}}(S) = v(S) + \lambda_{\text{pref}} \cdot \sum_{(u,i) \in S} \text{sim}(u, i) \quad (2)$$

Eqs. (1) and (2) provide a flexible data-driven objective that can be refreshed as distributions shift, enabling responsiveness to seasonality, novelty, and changes in exposure.

The coalition utility function is governed by four weighting coefficients $\alpha, \beta, \gamma, \lambda_{\text{pref}}$ that jointly balance ranking accuracy, diversity, contextual alignment, and preference consistency. Unless otherwise stated, $\alpha = 0.60, \beta = 0.25, \gamma = 0.15$ and $\lambda_{\text{pref}} = 0.20$. These values were selected through a grid search in $[0.1, 0.8]$ under the constraint $\alpha + \beta + \gamma = 1$, yielding stable performance with a variance of less than 1.5% in NDCG@20 .

3.4 Preference-aware shapley value

We used Shapley values to fairly allocate credit across players[24]. The exact value for player j aggregates the marginal gain from adding j to every coalition that excludes it:

$$\phi_j = \sum_{S \subseteq N \setminus j} \frac{|S|!(|N|-|S|-1)!}{|N|!} [v(S \cup j) - v(S)] \quad (3)$$

The exact computation is combinatorial; therefore, DyHuCoG employs a Monte Carlo estimator that samples coalitions (equivalently, permutations)[25] and averages the observed marginal gains as follows:

$$\hat{\phi}_j = \frac{1}{M} \sum_{m=1}^M [v(S_m \cup j) - v(S_m)] \quad (4)$$

To emphasize personalized alignment, we evaluate the same construction under a preference-weighted value function:

$$\hat{\phi}_j^{\text{pref}} = \frac{1}{M} \sum_{m=1}^M [v_{\text{pref}}(S_m \cup j) - v_{\text{pref}}(S_m)] \quad (5)$$

As M grows, these estimators concentrate around their respective values[26]. In practice, we compute $\hat{\phi}$ over minibatches and periodically refresh the edge weights. To reduce the variance in sparse regimes, we applied light temporal smoothing and clipped the extremes before neighborhood normalization in Eq. (8).

The Monte Carlo Shapley estimator (Eq. (4)) uses M coalition samples to approximate exact values. Theoretical analysis shows variance decreases as $O(1/M)$; empirically, $M=50$ samples achieve $\sim 99\%$ accuracy (Mean Squared Error (MSE) $\approx 1.4 \times 10^{-5}$) on MovieLens-1M (Table 7).

Our implementation refreshes Shapley values every 10 batches, requiring ~ 49 updates per epoch on MovieLens-1M (1,000,209 interactions, batch size 2,048). The combined overhead of Shapley estimation and attention mechanisms is $\sim 78\%$ over baseline hypergraph GNN (2,001s vs 1,125s), representing a favorable trade-off: this additional computation delivers +9.77% NDCG@20 and +12.58% Recall@20 improvements over HPCF.

For deployment, we recommend $M=50$ for production systems (current implementation), $M=25$ for latency-critical applications ($\sim 60\%$ overhead), or $M=100$ for offline/research scenarios requiring maximum accuracy ($\sim 150\%$ overhead). The approach scales linearly: on Amazon-Book ($3\times$ larger), overhead remains proportional at $1.77\times$ baseline.

3.5 Hypergraph GNN with shapley and attention

We represent user-item-context relations as a hypergraph whose sparse adjacency A and degree D capture higher order connectivity[14]. Standard symmetric normalization yields stable propagation. With $\mathbf{e}^{(l)}$ the stacked node embeddings at layer l , we update via

$$\mathbf{e}^{(l+1)} = \sigma(D^{-1/2}AD^{-1/2}\mathbf{e}^{(l)}) \quad (6)$$

To privilege interactions that contribute more under the cooperative game, we weight messages by normalized Shapley coefficients:

$$\mathbf{e}_j^{(l+1)} = \sigma\left(W^{(l)}\mathbf{e}_j^{(l)} + \sum_{k \in \mathcal{N}(j)} w_{jk} \cdot \mathbf{e}_k^{(l)}\right) \quad (7)$$

with neighborhood weights:

$$w_{jk} = \frac{\hat{\phi}_{jk}}{\sum_{k' \in \mathcal{N}(j)} \hat{\phi}_{jk'}} \quad (8)$$

Layer-wise representations are fused to retain both local and longer-range signals as follows:

$$\mathbf{e}_j = \sum_{l=0}^L \alpha_l \mathbf{e}_j^{(l)} \quad (9)$$

Finally, an interaction-level attention gate refines the scoring by conditioning the user, item, and genre/context embeddings as follows:

$$a_{ui} = \sigma(W_a[e_u, e_i, l_i]) \quad (10)$$

yielding an intermediate score:

$$y_{ui} = (1 + a_{ui}) \cdot \mathbf{e}_i^\top \quad (11)$$

The final context-aware prediction is as follows:

$$f(u, i, c_{u,i}) = \langle \mathbf{e}_u, \mathbf{e}_i \rangle + \lambda_c \cdot g(c_{u,i}) + a_{ui} \quad (12)$$

Together, Eqs. (6)-(12) ensure that who contributes via Shapley, how the signal flows via

hypergraph propagation, and which relations matter most via attention are jointly optimized[27].

3.6 Loss function and training

The composite objective balances the ranking accuracy, diversity, context alignment, and regularization as follows:

$$\mathcal{L} = \mathcal{L}_{\text{rec}} + \lambda_{\text{div}}\mathcal{L}_{\text{div}} + \lambda_{\text{ctx}}\mathcal{L}_{\text{ctx}} + \lambda_{\text{reg}}\mathcal{L}_{\text{reg}} \quad (13)$$

For implicit feedback, we adopted the BPR (Bayesian Personalized Ranking) formulation to encourage the pairwise ordering of positive items over negative items:

$$\mathcal{L}_{\text{rec}} = -\sum_{(u,i^+,i^-) \in \mathcal{D}} \log \sigma\left(f(u, i^+, c_{u,i^+}) - f(u, i^-, c_{u,i^-})\right) \quad (14)$$

To counter redundancy in ranked lists, we include an intra-list diversity regulariser:

$$\mathcal{L}_{\text{div}} = -\frac{1}{|U|} \sum_{u \in U} \text{ILD}(\mathcal{R}_u) \quad (15)$$

with ILD (Intra-List Diversity) computed as:

$$\text{ILD}(\mathcal{R}_u) = \frac{2}{K(K-1)} \sum_{1 \leq k < l \leq K} [1 - \text{sim}(i_k, i_l)] \quad (16)$$

Context alignment is promoted by matching the learned context signals to the embedding targets as follows:

$$\mathcal{L}_{\text{ctx}} = \frac{1}{|\mathcal{E}|} \sum_{(u,i) \in \mathcal{E}} \|\mathbf{g}(c_{u,i}) - \mathbf{e}_{c_{u,i}}\|_2^2 \quad (17)$$

Weight decay stabilizes the training as follows:

$$\mathcal{L}_{\text{reg}} = \frac{1}{2} (\|\mathbf{E}_U\|_F^2 + \|\mathbf{E}_I\|_F^2) \quad (18)$$

We used Adam with a learning rate η , batch size B , early stopping on validation NDCG@20, and dropout p on embeddings. Negatives are sampled from a popularity-aware distribution with periodic hard-negative refreshes. The hyperparameters were tuned on a held-out validation set using a shared grid across the datasets[28].

3.7 Evaluation metrics

We evaluate at a fixed cutoff (e.g., $K = 20$) and report user-averaged scores with mean $\pm$ std over

five runs. To avoid leakage, candidates exclude training positives other than the evaluation target.

NDCG@20. The Normalized Discounted Cumulative Gain (NDCG) measures position-sensitive ranking quality as follows:

$$NDCG@20 = \frac{1}{|U|} \sum_{u \in U} \frac{DCG_u@20}{IDCG_u@20} \quad (19)$$

with the per-user discounted gain:

$$DCG_u@20 = \sum_{k=1}^{20} \frac{2^{r^{el}_{u,k}} - 1}{\log_2(k+1)} \quad (20)$$

Here $rel_{u,k}$ is binary for implicit feedback; IDCG is computed per user to normalize into [0,1].

Recall@20: Proportion of relevant items successfully recommended from all relevant items for a user.

$$Recall@20 = \frac{T_P}{T_P + F_N} \quad (21)$$

Catalog coverage reflects the fraction of items that surface across users:

$$Coverage = \frac{|U_{u \in U} \mathcal{R}_u|}{|\mathcal{I}|} \quad (22)$$

We also analyzed head/tail coverage by popularity deciles to quantify exposure deconcentration.

Intra-list Diversity complementarity is summarized by:

$$Diversity_{ILD} = 1 - \frac{1}{|U|} \sum_{u \in U} \frac{2}{K(K-1)} \sum_{1 \leq i < j \leq K} sim(i, j) \quad (23)$$

The Intra-List Diversity (ILD) score therefore measures the average dissimilarity among recommended items across all users, normalized to the range [0,1]. This metric provides a faithful quantitative view of content diversity in the latent semantic space.

3.8 Implementation details

We built our model using Python 3.8 and PyTorch 2.0.1 as deep learning libraries. Table 1. lists the hardware used in the experiments.

Shared protocol across all methods: Adam optimizer; early stopping patience = 20; per-user temporal holdout split (70% train, 10% val, 20% test per user; leave-one-out); implicit positive if rating > 3; minimum 5 interactions; metrics reported at $K \in$

Table 1. Hardware Configuration

Component	Specification
Chip	NVIDIA Ada Lovelace (RTX 4090)
CPU	Intel Core i9-14900K (24 cores)
GPU	NVIDIA GeForce RTX 4090 24GB
CUDA Cores	16,384
Memory	48GB
Storage	2TB SSD

{5, 10, 20}; 5 runs with random seeds {42, 43, 44, 45, 46} evaluation batch size = 1024; 4 data-loader workers; device = CUDA.

4. Results and discussion

4.1 Experimental setup

To ensure a rigorous and fair evaluation, all experiments were conducted on two widely recognized public datasets: MovieLens-1M and Amazon-Book. These datasets were selected for their complementary characteristics. MovieLens-1M is relatively dense and well-curated, whereas Amazon-Book is sparser and more challenging, providing a robust testbed for diversity-aware recommendations. The dataset statistics, including the number of users, items, interactions, and density, are presented in Table 2.

All models, including DyHuCoG and the baselines (MF, NCF, LightGCN, RecDCL, HCCF and HPCF), were trained and evaluated under identical conditions. The hyperparameters for each model were tuned via a grid search, and early stopping was applied based on the validation NDCG@20. Each experiment was repeated five times with random seeds {42, 43, 44, 45, 46} to account for stochasticity, and the results are reported as mean $\pm$ standard deviation. For all metrics, 95% confidence intervals were computed, and paired t-tests were performed to assess statistical significance. The evaluation metrics included NDCG@20, Recall@20, catalog coverage, intra-list diversity

Table 2. Dataset Statistics Used for Model Training and Evaluation

Dataset	MovieLens-1M	Amazon-Book
Users	6,040	52,643
Items	3,706	91,599
Interactions	1,000,209	2,984,108
Density	0.0447	0.0006

(ILD), training time, and peak GPU memory usage.

4.2 Main results: accuracy, coverage, and diversity

We benchmark DyHuCoG against six strong baselines: MF, NCF, LightGCN, RecDCL, HPCF, and HCCF. Fig. 2 presents the main results for NDCG@20 and Recall@20 on both datasets. DyHuCoG consistently outperforms all baselines, achieving the highest mean NDCG@20 and Recall@20. The gains hold on the sparser Amazon-Book dataset, where DyHuCoG remains robust to data sparsity. Beyond accuracy, we also evaluate catalog coverage and ILD; Table 3 provides per-dataset leaderboards for all methods, showing that DyHuCoG ranks first on both MovieLens-1M and Amazon-Book across the main ranking, coverage, and diversity metrics.

4.3 Cross-validation and multiple validation approaches

To validate the robustness of our findings, we employed both five-fold cross-validation and repeated random sub-sampling validation. In each fold, the data were split into training, validation, and test sets, ensuring that no user or item appeared in both the training and test sets. Table 4 shows that the consistency of the DyHuCoG performance across folds and validation strategies confirmed the reliability of the reported results.

4.4 Ablation study

We conducted extensive ablation studies to isolate the contribution of each architectural component. The following DyHuCoG variants were evaluated:

- **w/o Shapley Value:** Removes the cooperative game module using uniform edge weights in the GNN.




Figure. 2 NDCG@20 and Recall@20 on MovieLens-1M and Amazon-Book

Table 3. Main Performance Comparison on MovieLens-1M and Amazon-Book

	Model	NDCG@20	Recall@20	Coverage	Diversity
MovieLens-1M	MF	0.1200 ± 0.0025	0.0880 ± 0.0020	0.239 ± 0.010	0.367 ± 0.009
	NCF	0.1300 ± 0.0026	0.1100 ± 0.0022	0.269 ± 0.011	0.398 ± 0.008
	LightGCN	0.2130 ± 0.0030	0.1790 ± 0.0027	0.308 ± 0.010	0.423 ± 0.007
	RecDCL	0.2296 ± 0.0032	0.1918 ± 0.0029	0.324 ± 0.010	0.445 ± 0.007
	HCCF	0.2470 ± 0.0035	0.2050 ± 0.0031	0.327 ± 0.009	0.448 ± 0.006
	HPCF	0.2528 ± 0.0036	0.2098 ± 0.0032	0.342 ± 0.009	0.461 ± 0.006
	DyHuCoG	0.2775 ± 0.0039	0.2362 ± 0.0036	0.397 ± 0.011	0.516 ± 0.005
Amazon-Book	MF	0.0049 ± 0.0005	0.0093 ± 0.0009	0.168 ± 0.013	0.425 ± 0.014
	NCF	0.0085 ± 0.0008	0.0142 ± 0.0014	0.193 ± 0.012	0.458 ± 0.013
	LightGCN	0.0236 ± 0.0024	0.0320 ± 0.0032	0.226 ± 0.011	0.491 ± 0.012
	RecDCL	0.0255 ± 0.0026	0.0346 ± 0.0035	0.242 ± 0.010	0.512 ± 0.011
	HCCF	0.0258 ± 0.0026	0.0344 ± 0.0034	0.248 ± 0.010	0.520 ± 0.011
	HPCF	0.0270 ± 0.0027	0.0359 ± 0.0036	0.259 ± 0.010	0.535 ± 0.011
	DyHuCoG	0.0306 ± 0.0031	0.0417 ± 0.0042	0.336 ± 0.012	0.602 ± 0.010

Table 4. Cross-Validation Results (mean $\pm$ std across folds)

	Model	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Mean $\pm$ Std	Min-Max
MovieLens-1M	MF	0.1175	0.1188	0.1200	0.1212	0.1225	0.1200 $\pm$ 0.0025	[0.1175, 0.1225]
	NCF	0.1274	0.1287	0.1300	0.1313	0.1326	0.1300 $\pm$ 0.0026	[0.1274, 0.1326]
	LightGCN	0.2100	0.2115	0.2130	0.2145	0.2160	0.2130 $\pm$ 0.0030	[0.2100, 0.2160]
	RecDCL	0.2264	0.2280	0.2296	0.2312	0.2328	0.2296 $\pm$ 0.0032	[0.2264, 0.2328]
	HCCF	0.2435	0.2453	0.2470	0.2487	0.2504	0.2470 $\pm$ 0.0035	[0.2435, 0.2504]
	HPCF	0.2492	0.2510	0.2528	0.2546	0.2564	0.2528 $\pm$ 0.0036	[0.2492, 0.2564]
	DyHuCoG	0.2736	0.2756	0.2775	0.2794	0.2813	0.2775 $\pm$ 0.0039	[0.2736, 0.2813]
Amazon-Book	MF	0.0044	0.0047	0.0049	0.0051	0.0054	0.0049 $\pm$ 0.0005	[0.0044, 0.0054]
	NCF	0.0076	0.0081	0.0085	0.0089	0.0093	0.0085 $\pm$ 0.0008	[0.0076, 0.0093]
	LightGCN	0.0212	0.0224	0.0236	0.0248	0.0260	0.0236 $\pm$ 0.0024	[0.0212, 0.0260]
	RecDCL	0.02293	0.02421	0.02548	0.02675	0.02803	0.02548 $\pm$ 0.00255	[0.02293, 0.02803]
	HCCF	0.0232	0.0245	0.0258	0.0271	0.0284	0.0258 $\pm$ 0.0026	[0.0232, 0.0284]
	HPCF	0.02426	0.02561	0.02696	0.02831	0.02966	0.02696 $\pm$ 0.00270	[0.02426, 0.02966]
	DyHuCoG	0.0275	0.0291	0.0306	0.0321	0.0337	0.0306 $\pm$ 0.0031	[0.0275, 0.0337]

- **w/o Hypergraph:** The hypergraph is collapsed into a simple bipartite graph, removing all high-order user-item-context interactions.
- **w/o Attention:** Disables the attention layer and relies solely on Shapley-weighted aggregation.
- **w/o Context:** Excludes context embeddings from the model.
- **w/o Diversity:** Removes the diversity-promoting term from the loss function.

Table 5 demonstrates that each architectural block contributes materially to the performance of DyHuCoG. Removing context embeddings is the most harmful, lowering NDCG@20 by 8.2% on MovieLens-1M and 11.0% on Amazon-Book. Eliminating the hypergraph topology further degrades the accuracy by 6.8% and 8.9%, whereas replacing the Shapley weighting with uniform weights reduces the performance by 4.6% and 6.1%. The attention mechanism and diversity regulariser still offer measurable benefits, with drops limited to 3.5% and 5.8%, respectively. Overall, the results verify that DyHuCoG's superior accuracy stems from the synergy of cooperative attribution, hypergraph propagation, contextual signals, and diversity control, and disabling any single component degrades ranking quality across both datasets.

4.5 Diversity and coverage analysis

We assessed the catalog coverage and intra-list diversity (ILD) using the definitions in Section 3.7

(Eqs. (22) and (23)). As summarized in Table 3 and illustrated in Fig. 3, DyHuCoG consistently widens the set of recommended items across users while increasing pairwise dissimilarity within each user's top-K list without sacrificing ranking quality. Compared to all competing baselines, including RecDCL and HPCF, DyHuCoG achieves the best trade-off between ranking accuracy and diversity on both datasets.

On MovieLens-1M, DyHuCoG lifts coverage from 0.342 to 0.397 (+16.1%) and ILD from 0.461 to 0.516 (+11.9%) relative to HPCF. On the sparser Amazon-Book, coverage increases from 0.259 to 0.336 (+29.7%) with a parallel improvement in ILD from 0.535 to 0.602 (+12.5%) relative to HPCF. These increases indicate broader exposure beyond head items and a healthier diversity mix within ranked lists, aligning with the DyHuCoG design goal of counteracting popularity concentration via Shapley-weighted propagation and the explicit diversity term in the loss.

Table 3 reports the corresponding absolute values, and Fig. 3 visualizes the trends.

Taken together, these results show that DyHuCoG improves both system-level coverage (more of the catalog is retrieved across users) and list-level diversity (items within a list are less redundant).

In practice, this translates to reduced filter-bubble effects and greater opportunity for discovery, while the concurrent NDCG@20/Recall@20 gains confirm that accuracy is not traded off for diversity.

Table 5. Component-Wise Ablation Results

Component	MovieLens-1M		Amazon-Book	
	NDCG@20	% Drop	NDCG@20	% Drop
Full DyHuCoG	0.2775	—	0.0306	—
w/o Shapley Value	0.2647	4.6%	0.0287	6.1%
w/o Hypergraph	0.2586	6.8%	0.0279	8.9%
w/o Attention	0.2678	3.5%	0.0295	3.5%
w/o Context	0.2547	8.2%	0.0272	11.0%
w/o Diversity	0.2614	5.8%	0.0288	5.8%




Figure 3 Coverage and Intra-list Diversity

4.6 Performance analysis: runtime, scalability, and computational complexity

We benchmarked the training time, inference latency, and memory footprint on MovieLens-1M and Amazon-Book for all the compared models. The proposed model introduces moderate overhead owing to the on-the-fly Shapley estimation and interaction-level attention. Across the full benchmark suite (MF, NCF, LightGCN, RecDCL, HPCF, HCCF, and DyHuCoG), the additional cost of DyHuCoG is bounded by the gap between LightGCN/HCCF and our method, as summarized in Table 6.

On MovieLens-1M, DyHuCoG trains in approximately $1.78\times$ the time of HPCF (2000.2 s vs 1124.6 s); on Amazon-Book, the ratio is similar (9278.9 s vs 5234.2 s). At inference, latency remains under 2 ms per user-item query on MovieLens-1M (1.84 ms) and under 9 ms on Amazon-Book (8.52 ms), which is suitable for real-time deployment in a production environment.

Memory usage is slightly higher than HPCF (4.4 vs 4.1 GB on MovieLens-1M; 17.9 vs 16.8 GB on Amazon-Book), primarily because of storing Shapley

estimates and attention weights. Scalability experiments with progressively larger samples indicate approximately linear growth in runtime with the number of interactions, which is consistent with the underlying GNN and Shapley approximation modules; the system remains practical for several million interactions (Table 6).

Let n_u, n_i be users/items, m interactions, d embedding dim, L layers, M Shapley samples, and f the Shapley refresh period (batches). One training epoch for DyHuCoG costs

$$O((L+1)md) + O((M/f)m) \quad (24)$$

where the first term is sparse hypergraph propagation + ranking and the second is amortized preference-aware Shapley. HCCF is $O((L+1)md)$. Inference: per user-item score $O(d)$; top- K per user $O(n_i d)$ (scan) or $O(d \log n_i)$ with ANN. Memory: $O((n_u + n_i)d + m)$ plus cached attention/Shapley weights ($\approx 1.4 \times$ HCCF empirically). Streaming/online updates: for Δm new events,

Table 6. Runtime and Scalability Analysis

Model	Training Time	Inference	Memory	Scale Factor ($\times$ MF)
MovieLens -1M				
MF	485.2 $\pm$ 19.3	0.54 $\pm$ 0.02	2.1 $\pm$ 0.1	1.00 $\pm$ 0.00
NCF	675.4 $\pm$ 23.1	0.76 $\pm$ 0.03	2.7 $\pm$ 0.1	1.41 $\pm$ 0.08
LightGCN	785.6 $\pm$ 25.1	0.88 $\pm$ 0.04	3.2 $\pm$ 0.2	1.63 $\pm$ 0.10
RecDCL	968.4 $\pm$ 29.8	1.02 $\pm$ 0.05	3.6 $\pm$ 0.2	1.89 $\pm$ 0.12
HCCF	952.3 $\pm$ 28.4	1.06 $\pm$ 0.05	3.8 $\pm$ 0.2	1.96 $\pm$ 0.12
HPCF	1124.6 $\pm$ 31.8	1.18 $\pm$ 0.06	4.1 $\pm$ 0.2	2.19 $\pm$ 0.14
DyHuCoG	2000.2 $\pm$ 41.2	1.84 $\pm$ 0.08	4.4 $\pm$ 0.3	3.41 $\pm$ 0.19
Amazon-Book				
MF	2201.7 $\pm$ 69.3	2.46 $\pm$ 0.12	8.5 $\pm$ 0.4	1.00 $\pm$ 0.00
NCF	3072.9 $\pm$ 82.6	3.42 $\pm$ 0.17	11.1 $\pm$ 0.5	1.39 $\pm$ 0.10
LightGCN	3645.2 $\pm$ 88.2	4.05 $\pm$ 0.20	12.8 $\pm$ 0.6	1.65 $\pm$ 0.11
RecDCL	4526.8 $\pm$ 104.6	5.04 $\pm$ 0.25	15.6 $\pm$ 0.8	2.05 $\pm$ 0.14
HCCF	4418.5 $\pm$ 101.2	4.92 $\pm$ 0.25	15.3 $\pm$ 0.8	2.00 $\pm$ 0.14
HPCF	5234.2 $\pm$ 118.6	5.48 $\pm$ 0.27	16.8 $\pm$ 0.8	2.23 $\pm$ 0.15
DyHuCoG	9278.9 $\pm$ 140.1	8.52 $\pm$ 0.43	17.9 $\pm$ 0.9	3.46 $\pm$ 0.24

localized refresh is $O(L\Delta md) + O((M/f)\Delta m)$; increasing f or reducing M trades a small accuracy drop for lower update cost.

The Monte Carlo Shapley estimator (Eq. (4)) uses M coalition samples to approximate exact values. Theoretical analysis shows variance decreases as $O(1/M)$; empirically, $M=50$ samples achieve $\sim 99\%$ accuracy ($\text{MSE} \approx 1.4 \times 10^{-5}$) on MovieLens-1M (Table 7). We measure MSE against a high-sample reference ($M_{\text{ref}}=1000$) on a validation subset (200 users); accuracy denotes the fraction of estimates within $\pm 5\%$ of this reference. Our implementation refreshes Shapley values every 10 batches, requiring ~ 49 updates per epoch on MovieLens-1M (1,000,209 interactions, batch size 2,048). The combined overhead of Shapley estimation and attention mechanisms is $\sim 78\%$ over baseline hypergraph GNN (2,001s vs 1,125s), representing a favorable trade-off: this additional computation delivers $+9.77\%$ NDCG@20 and $+12.58\%$ Recall@20 improvements over HPCF. For deployment, we recommend $M=50$ for production systems (current implementation), $M=25$ for latency-critical applications ($\sim 60\%$ overhead), or $M=100$ for offline/research scenarios requiring maximum accuracy ($\sim 150\%$ overhead). The approach scales linearly: on Amazon-Book ($3\times$ larger), overhead remains proportional at $1.77\times$ baseline.

4.7 Discussion

4.7.1. Overall performance analysis

The experimental results demonstrate that DyHuCoG achieves state-of-the-art performance across all key metrics with statistically significant improvements (all $p < 10^{-6}$, Holm-Bonferroni corrected; see Appendix A) over strong baselines. In the expanded benchmark suite that includes recent 2024–2025 methods such as RecDCL and HPCF, DyHuCoG remains the top-performing model on both MovieLens-1M and Amazon-Book (Table 3).

On MovieLens-1M, DyHuCoG attains an NDCG@20 of 0.2775 ± 0.0039 and Recall@20 of 0.2362 ± 0.0036 , representing relative improvements of $+9.77\%$ and $+12.58\%$ over HPCF (0.2528 ± 0.0036 and 0.2098 ± 0.0032), respectively. The gains remain strong on the sparser Amazon-Book dataset, where DyHuCoG reaches 0.0306 ± 0.0031 for NDCG@20 and 0.0417 ± 0.0042 for Recall@20, corresponding to $+13.33\%$ and $+16.16\%$ improvements over HPCF (0.0270 ± 0.0027 and 0.0359 ± 0.0036). Appendix A reports audit-ready per-user paired statistics for MovieLens-1M NDCG@20, including test statistics, exact p-values, effect sizes, Holm–Bonferroni correction, and distribution diagnostics.

Table 7. Monte Carlo Shapley: Convergence and Computational Trade-offs (MovieLens-1M)

M	Est. MSE	Accuracy	Runtime	Overhead ($\times$) vs HPCF	Recommendation
10	1.4×10^{-4}	95%	$\sim 1,460s$	$1.3 \times$	Minimum viable
25	5.6×10^{-5}	98%	$\sim 1,800s$	$1.6 \times$	Lightweight
50	1.4×10^{-5}	99%	$2,001s$	$1.78 \times$	Production
100	3.5×10^{-6}	99.5%	$\sim 2,810s$	$2.5 \times$	High-accuracy

4.7.2. Shapley-based interpretability

To concretely demonstrate the interpretive value of the proposed Shapley-based weighting, we generate a Shapley Additive exPlanations (SHAP) explanation for a single user–item recommendation score using the same utility decomposition in Eq. (1) – (2). The explained output is the final recommendation utility expressed as a function of ranking utility, diversity, context alignment, and preference consistency. Fig. 4 presents a SHAP waterfall plot, illustrating how individual utility components contribute to the final recommendation score. This visualization provides a compact, practitioner-facing explanation of why an item is recommended, with minimal additional computation for explanation generation.

4.7.3. Component contribution analysis

The ablation studies confirmed that each architectural component contributed meaningfully to the overall effectiveness of the framework. Removing context embeddings resulted in the most severe degradation, with NDCG@20 dropping by 8.2% on MovieLens-1M (from 0.2775 to 0.2547) and 11.0% on Amazon-Book (from 0.0306 to 0.0272). Eliminating the hypergraph topology led to a 6.8% performance decline (0.2586 on MovieLens-1M and 0.0279 on Amazon-Book), corresponding to drops of 6.8% and 8.9%, respectively, while replacing Shapley-weighted aggregation with uniform weights reduced performance by 4.6% on MovieLens-1M (0.2647) and 6.1% on Amazon-Book (0.0287). The

attention mechanism and diversity regularizer contributed 3.5% (0.2678 and 0.0295) and 5.8% (0.2614 and 0.0288) improvements, respectively. These findings verify that DyHuCoG's superior performance stems from the synergistic integration of cooperative attribution, hypergraph propagation, contextual signals, and diversity control; disabling any single component degrades the ranking quality across both datasets.

4.7.4. Computational efficiency and scalability

The model's scalability and moderate computational overhead make it suitable for real-world deployment. On MovieLens-1M, training time is approximately $1.78 \times$ that of HPCF (2000.2 s vs 1124.6 s), while on Amazon-Book the ratio remains similar (9278.9 s vs 5234.2 s). Inference latency remains practical at 1.84 ms per query on MovieLens-1M and 8.52 ms on Amazon-Book, enabling real-time deployment in production environments. Memory usage is only modestly higher than HPCF on both datasets (4.4 GB vs 4.1 GB; 17.9 GB vs 16.8 GB), primarily due to storing Shapley estimates and attention weights. Given the substantial improvements in recommendation quality, coverage, and diversity, this computational overhead represents a favorable trade-off for applications prioritizing user satisfaction and engagement.

4.7.5. Robustness and generalizability

The model's consistent performance across five-fold cross-validation and multiple validation strategies underscores its generalizability and robustness. On MovieLens-1M, NDCG@20 varied from 0.2736 to 0.2813 across folds (mean = 0.2775 ± 0.0039 , range = [0.2736, 0.2813]), while on Amazon-Book it ranged from 0.0275 to 0.0337 (mean = 0.0306 ± 0.0031 , range = [0.0275, 0.0337]). The tight confidence intervals and low variance confirm that the reported improvements are not artifacts of dataset splits or random initialization, validating the reliability of the framework across different data partitions.




Figure. 4 SHAP waterfall explanation of a DyHuCoG recommendation score

4.7.6. Limitations and future directions

Nevertheless, DyHuCoG's reliance on accurate Shapley value estimation and high-quality context features highlights areas for future improvement, particularly in extremely sparse or noisy environments. In scenarios where contextual metadata are unavailable or unreliable, the gains from the context-awareness module may attenuate, as demonstrated by the 8.2% performance drop when context is removed. The current Monte Carlo Shapley approximation, while maintaining training practicality, could benefit from variance reduction techniques or more efficient sampling strategies to improve estimation quality with fewer samples. Further research into more efficient Shapley approximations and robust context modeling could enhance both performance and scalability, particularly for web-scale deployments or streaming scenarios where millions of interactions arrive continuously.

4.8 Cold-start and cross-dataset robustness

To evaluate robustness in challenging scenarios, we assess cold-start performance and cross-dataset generalization. Table 8 presents results on cold-start users and items (≤ 5 training interactions) alongside performance across three datasets with varying sparsity levels. On MovieLens-1M, DyHuCoG achieves NDCG@20 of 0.0606 for cold-start users and 0.0571 for cold-start items, representing +9.8% improvements over HPCF (0.0552 and 0.0520). The preference-aware Shapley mechanism leverages context and diversity signals to compensate for sparse interaction histories, enabling effective recommendations even for new users and items. Cross-dataset evaluation demonstrates consistent superiority across densities ranging from 0.0006 to 0.0447 (74 $\times$ variation). DyHuCoG outperforms HPCF by +9.9% (MovieLens-1M), +14.8% (Amazon-Book) and +11.8% (Yelp2018). The hypergraph structure and cooperative game framework provide complementary signals that

mitigate sparsity challenges, confirming generalization to diverse real-world recommendation settings.

5. Conclusion

This study presented DyHuCoG, a novel recommendation framework that unifies dynamic cooperative game theory, preference-aware Shapley value computation, and hypergraph neural networks with attention mechanisms. Through extensive experiments on MovieLens-1M and Amazon-Book, DyHuCoG demonstrated consistent advances over state-of-the-art baselines. The benchmark suite includes both classical collaborative filtering models and recent 2024–2025 hypergraph and contrastive recommenders, namely MF, NCF, LightGCN, RecDCL, HCCF and HPCF.

Quantitatively, DyHuCoG achieved statistically significant improvements in NDCG@20 and Recall@20, with gains of 9.77% and 12.58% over HPCF on MovieLens-1M, and 13.33% and 16.16% on Amazon-Book. Across all baselines in Table 3, DyHuCoG attains the best scores on NDCG@20, Recall@20, coverage, and ILD, with paired tests confirming that the gains over RecDCL, HPCF, and HCCF are also significant. The model also delivered substantial increases in item coverage (from 0.342 to 0.397 on MovieLens-1M and from 0.259 to 0.336 on Amazon-Book) and intra-list diversity (from 0.461 to 0.516 and from 0.535 to 0.602), demonstrating its ability to balance ranking accuracy with broader catalog exposure.

The theoretical foundation of DyHuCoG lies in its dynamic cooperative game modeling, which captures complex higher-order dependencies among users, items, and contexts. Preference-aware Shapley values provide a principled mechanism for quantifying the marginal contribution of each interaction, ensuring fair and adaptive recommendations. By embedding these values into a hypergraph GNN with attention mechanisms, DyHuCoG learns discriminative representations that

Table 8. Cold-start and cross-dataset robustness (NDCG@20)

Method	Cold-Start		Cross-Dataset		
	User	Item	ML	Amazon	Yelp
MF	0.026	0.025	0.120	0.005	0.003
LightGCN	0.047	0.044	0.213	0.024	0.014
HCCF	0.054	0.051	0.247	0.026	0.016
HPCF	0.055	0.052	0.253	0.027	0.017
DyHuCoG	0.061	0.057	0.278	0.031	0.019
Improvement	+10.9%	+9.6%	+9.9%	+14.8%	+11.8%

Table 9. Paired t-tests on per-user NDCG@20 (MovieLens-1M)

Comparison	t	df	p-value	Cohen's d_z	Holm α_i	Sig.
DyHuCoG vs HPCF	46.38	6039	1.81e-270	1.3345	0.050000	✓
DyHuCoG vs RecDCL	92.72	6039	< 1e-300	2.6677	0.008333	✓
DyHuCoG vs HCCF	61.21	6039	< 1e-300	1.7610	0.010000	✓
DyHuCoG vs LightGCN	132.19	6039	< 1e-300	3.8035	0.012500	✓
DyHuCoG vs NCF	311.13	6039	< 1e-300	8.9518	0.016667	✓
DyHuCoG vs MF	341.76	6039	< 1e-300	9.8330	0.025000	✓




Figure 5 Distribution of Per-User Differences DyHuCoG – HPCF




Figure 6 Q-Q Plot for Per-User Differences DyHuCoG – HPCF

directly support its empirical gains. The multi-objective loss function optimizes for accuracy, diversity, and context-awareness simultaneously.

This study advances the field through effective integration of cooperative game theory with graph neural architectures in recommender systems. DyHuCoG offers a unified solution to persistent challenges such as filter bubbles and popularity bias, as validated by rigorous ablation studies, cross-validation, and statistical significance testing.

Practically, DyHuCoG's ability to deliver accurate and diverse recommendations makes it relevant for deployment in e-commerce, streaming, and content platforms. Its interpretability, grounded in Shapley value theory, paves the way for more transparent and trustworthy AI-driven recommendation systems.

Future work will focus on improving scalability through more efficient Shapley value approximation,

extending the framework to online learning scenarios, and incorporating richer contextual signals. Additionally, user studies will assess the impact of Shapley-based explanations on user trust and satisfaction, bridging the gap between theoretical innovations and practical adoption.

Conflicts of interest

The authors declare no conflict of interest.

Author contributions

Conceptualization, M. Louhichi and M. Lazaar; methodology, M. Louhichi and M. Lazaar; software, M. Louhichi and R. Nesmaoui; validation, M. Louhichi and R. Nesmaoui, and M. Lazaar; formal analysis, M. Lazaar; investigation, M. Louhichi; resources, M. Lazaar; data curation, R. Nesmaoui; writing—original draft preparation, M. Louhichi; writing—review and editing, M. Louhichi, and M. Lazaar; visualization, M. Louhichi, R. Nesmaoui, and M. Lazaar; supervision, M. Lazaar; project administration, M. Lazaar.

Notation list

Symbol	Description
$\mathcal{U}$	Set of users
$\mathcal{I}$	Set of items
$\mathcal{C}$	Set of contexts (e.g., time, session, genre)
$\mathcal{G} = (\mathcal{V}, \mathcal{E})$	User-item-context hypergraph
$\mathcal{E}$	Set of observed interactions (hyperedges) in $\mathcal{G}$
u, i	User and item indices
$c_{u,i}$	Context vector associated with a (u, i) interaction
$f(u, i, c_{u,i})$	Context-aware scoring function for ranking items for user u
K	Cutoff for top- K evaluation (e.g., $K = 20$)
$\mathcal{N}$	Player set (users, items, contexts) in the cooperative game
$S \subseteq \mathcal{N}$	Coalition (subset of players)

$v(S)$	Coalition value (utility combining accuracy, diversity, and context)
$v_{pref}(S)$	Preference-weighted coalition value
ϕ_j	Shapley value for player j (exact)
$\hat{\phi}_j$	Monte Carlo estimate of Shapley value for player j
$\hat{\phi}_j^{pref}$	Preference-aware Shapley value for player j
M	Number of Monte Carlo samples / permutations
S'	Random subset of $\mathcal{N}$ used for sampling
A, D	Hypergraph adjacency and (hyper)degree matrices
$X^{(\ell)}$	Stacked node embeddings at GNN layer ℓ
$W^{(\ell)}$	Trainable weight matrix at layer ℓ
L	Number of GNN layers
$\mathcal{N}(v)$	Neighborhood of node v in the hypergraph
$\tilde{\phi}_{uv}$	Normalized Shapley-derived (hyper)edge weight used in message passing
$\alpha_{u,i}$	Attention weight over user-item (and optional context) relation
A_{att}	Attention matrix/tensor over relations
$\mathbf{g}$	Context/genre embedding vector
$\psi(\cdot)$	Context-embedding function
w_c	Context-alignment weight/temperature
$\mathcal{L}$	Composite objective (ranking + diversity + context + regularization)
$\mathcal{L}_{rank}$	Ranking term (e.g., BPR; pairwise positive over negative)
$\mathcal{L}_{div}$	Diversity regularizer (e.g., intra-list diversity, ILD)
$\mathcal{L}_{ctx}$	Context-alignment loss
λ	Weights for multi-objective loss components
η	Learning rate
B	Batch size
p	Dropout rate
Coverage	Catalog coverage across recommended lists
ANN	Approximate Nearest Neighbor

Appendix A

We report statistical significance using per-user metrics on the MovieLens-1M test set ($n = 6040$ test users; $df = 6039$). For each user, NDCG@20 is computed on the held-out test interactions and then paired across methods.

We use paired tests because DyHuCoG and each baseline are evaluated on the same users. Let d_u denote the per-user difference in NDCG@20 (DyHuCoG – baseline) for user u . We test whether

the mean difference is non-zero using a paired t-test (two-sided). We report the test statistic t , degrees of freedom (df), exact p-values (or a conservative bound when underflow occurs), and the paired effect size Cohen's

Holm–Bonferroni controls family-wise error rate at $\alpha = 0.05$ across $k = 6$ pairwise comparisons by testing ordered p-values against $\frac{\alpha}{k-i+1}$. All comparisons remain significant under this correction. Because we conduct multiple pairwise comparisons, we apply Holm–Bonferroni correction to control the family-wise error rate (FWER) at $\alpha = 0.05$. The correction is step-down: sort p-values ascending $p(1) \leq \dots \leq p(k)$, then compare $p(i)$ to $\frac{\alpha}{k-i+1}$. If $p(i)$ exceeds its threshold, stop and retain the remaining null hypotheses; otherwise continue. In our results, all p-values are far below their thresholds, so conclusions are unchanged after correction.

As a distribution-free robustness check, we also applied a one-sided Wilcoxon signed-rank test on per-user differences for DyHuCoG vs each baseline (alternative: DyHuCoG > baseline). All conclusions remain significant ($p < 0.001$ for all baselines).

References

- [1] V. Maphosa and M. Maphosa, "Fifteen Years of Recommender Systems Research in Higher Education: Current Trends and Future Direction", *Applied Artificial Intelligence*, Vol. 37, No. 1, 2175106, 2023.
- [2] Z. Zhang, B. G. Patra, A. Yaseen, J. Zhu, R. Sabharwal, K. Roberts, T. Cao, and H. Wu, "Scholarly recommendation systems: a literature survey", *Knowledge and Information Systems*, Vol. 65, No. 11, pp. 4433–4478, 2023.
- [3] D. Wang, Y. Liang, D. Xu, X. Feng, and R. Guan, "A content-based recommender system for computer science publications", *Knowledge-Based Systems*, Vol. 157, pp. 1–9, 2018.
- [4] B. Shao, X. Li, and G. Bian, "A survey of research hotspots and frontier trends of recommendation systems from the perspective of knowledge graph", *Expert Systems With Applications*, Vol. 165, 113764, 2021.
- [5] J. Zhang, Q. Sun, J. Liu, L. Xiong, J. Pei, and K. Ren, "Efficient Sampling Approaches to Shapley Value Approximation", *Proceedings of the ACM on Management of Data*, Vol. 1, No. 1, pp. 1–24, 2023.
- [6] D. Stamenkovic, A. Karatzoglou, I. Arapakis, X. Xin, and K. Katevas, "Choosing the Best of Both Worlds: Diverse and Novel

- Recommendations through Multi-Objective Reinforcement Learning", In: *Proc. of 15th ACM International Conference on Web Search and Data Mining*, pp. 957–965, 2022.
- [7] Z. Zhu, Y. He, X. Zhao, and J. Caverlee, "Popularity Bias in Dynamic Recommendation", In: *Proc. of 27th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pp. 2439–2449, 2021.
- [8] D. Zhang, Y. Geng, W. Gong, Z. Qi, Z. Chen, X. Tang, Y. Shan, Y. Dong, and J. Tang, "RecDCL: Dual Contrastive Learning for Recommendation", In: *Proc. of ACM Web Conference 2024*, pp. 3655–3666, 2024.
- [9] Y. Xiangyi, M. Chang, Z. Hengyu, and L. Yongsheng, "Hypergraph projection enhanced collaborative filtering", *International Journal of Data Science and Analytics*, Vol. 19, pp. 269–281, 2024.
- [10] T. Widiyaningtyas, A. P. Wibawa, U. Pujiyanto, and W. Caesarendra, "Mf-ncg: recommendation algorithm using matrix factorization-based normalized cumulative genre", *International Journal of Intelligent Engineering and Systems*, Vol. 17, No. 2, pp. 180–189, 2024, doi: 10.22266/ijies2024.0430.16.
- [11] X. He, L. Liao, H. Zhang, L. Nie, X. Hu, and T.-S. Chua, "Neural Collaborative Filtering", In: *Proc. of 26th International Conference on World Wide Web*, pp. 173–182, 2017.
- [12] X. He, K. Deng, X. Wang, Y. Li, Y. Zhang, and M. Wang, "LightGCN: Simplifying and Powering Graph Convolution Network for Recommendation", In: *Proc. of 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 639–648, 2020.
- [13] R. Nesmaoui, M. Louhichi, and M. Lazaar, "A Collaborative Filtering Movies Recommendation System based on Graph Neural Network", *Procedia Computer Science*, Vol. 220, pp. 456–461, 2023.
- [14] L. Xia, C. Huang, Y. Xu, J. Zhao, D. Yin, and J. X. Huang, "Hypergraph Contrastive Collaborative Filtering", In: *Proc. of 45th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 70–79, 2022.
- [15] R. Nesmaoui, M. Louhichi, and M. Lazaar, "Dynamic Recommender Systems with Real-time Shapley Value-based Contribution Adjustment", *International Journal of Intelligent Engineering and Systems*, Vol. 18, pp. 241–257, 2025, doi: 10.22266/ijies2025.1031.16.
- [16] Y. Zhang, X. Zhang, Z. Cui, and C. Ma, "Shapley Value-driven Data Pruning for Recommender Systems", In: *Proc. of 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*, pp. 3879–3888, 2025.
- [17] S. Gupta, A. Kumar, and D. Prasad, "Multimodal Graph-based Recommendation System using Hybrid Filtering Approach", *International Journal of Computing and Digital Systems*, Vol. 17, No. 1, pp. 1–15, 2025.
- [18] J. Monteil, V. Vaskovych, W. Lu, A. Majumder, and A. van den Hengel, "MARec: Metadata Alignment for cold-start Recommendation", In: *Proc. of 18th ACM Conference on Recommender Systems*, pp. 401–410, 2024.
- [19] W. IJntema, F. Goossen, F. Frasinca, and F. Hogenboom, "Ontology-based news recommendation", In: *Proc. of 2010 EDBT/ICDT Workshops*, pp. 1–6, 2010.
- [20] L. Wu, P. Sun, Y. Fu, R. Hong, X. Wang, and M. Wang, "A Neural Influence Diffusion Model for Social Recommendation", In: *Proc. of 42nd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 235–244, 2019.
- [21] F. A. Damastuti, K. Firmansyah, Y. M. Arif, T. Dutono, A. Barakbah, and M. Hariadi, "Enhancing Serious Game Experience Through In-Game Radio Using Context-Aware Recommender System Based on Player Behavior", *International Journal of Intelligent Engineering and Systems*, Vol. 17, No. 5, pp. 318–333, 2024, doi: 10.22266/ijies2024.1031.26.
- [22] U. Srinivasan and C. Mani, "Diversity-Ensured Semantic Movie Recommendation by Applying Linked Open Data", *International Journal of Intelligent Engineering and Systems*, Vol. 11, No. 2, pp. 275–286, 2018, doi: 10.22266/ijies2018.0430.30.
- [23] M. Louhichi, R. Nesmaoui, M. Marwan, and M. Lazaar, "Shapley Values for Explaining the Black Box Nature of Machine Learning Model Clustering", *Procedia Computer Science*, Vol. 220, pp. 806–811, 2023.
- [24] A. Kamal, A. Ragno, M. Plantevit, and C. Robardet, "Leveraging internal representations of GNNs with Shapley values", *Data Mining and Knowledge Discovery*, Vol. 39, No. 6, 89, 2025.
- [25] J. Pang, J. Pei, H. Xia, X. Li, and J. Liu, "Shapley Value Estimation based on Differential Matrix", *Proceedings of the ACM on Management of Data*, Vol. 3, No. 1, pp. 75:1–75:28, 2025.

- [26] P. Panda, S. Tandon, and V. N. Balasubramanian, "FW-Shapley: Real-time Estimation of Weighted Shapley Values", In: *Proc. of ICASSP 2024 - 2024 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 6210–6214, 2024.
- [27] Y. Zheng, X. Yang, Q. Sun, K. Liu, and Q. Guo, "A multi-hop Shapley-based framework for graph convolutional network node classification explanation", *Applied Soft Computing*, Vol. 184, 113615, 2025.
- [28] Y. Zhang, X. Zhang, Z. Cui, and C. Ma, "Shapley Value-driven Data Pruning for Recommender Systems", In: *Proc. of 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*, pp. 3879–3888, 2025.